

We declare that we have completed this assignment in accordance with the UAB Academic Integrity Code and the UAB CS Honor Code We have read the UAB Academic Integrity Code and understand that any breach of the Code may result in severe penalties. We also declare that the following percentage distribution faithfully represents individual group member's contributions to the completion of the assignment.

Name	Overall Contribution (%)	Major work items completed by me	Signature/ Initials	Date
Vekataramana sai teja dasarathi	25	<p>Data Collection Implement Internet of Things sensors to monitor bin fill levels in real time.</p> <p>Use Google Maps API to geotag bin locations.</p> <p>Use K-Means clustering to categorise neighbourhoods according to garbage generation.</p>	VSD	Nov 15
Bhargavi Janapati	25	<p>Analyse experimental data to determine system efficacy.</p> <p>Prepare documentation outlining the techniques and findings.</p>	BJ	Nov 20

		Discuss possible future enhancements and research possibilities.		
Gurram Mourya	25	<p>Create dynamic collection schedules that prioritise high-waste zones.</p> <p>Provide user tips for better recycling methods.</p> <p>Create visualisations, such as heatmaps, to pinpoint waste hotspots.</p>	GM	Nov 25
Likith Kumar Tarala	25	<p>Use Decision Tree Regression to forecast garbage generation and optimise collection schedules.</p> <p>Implement the Apriori algorithm to identify trends in trash disposal behaviour.</p> <p>Prepare documentation detailing methodologies and findings.</p>	LKT	Nov 30

Smart Waste Management System Using Data Mining

Table of Contents

1. Preface

2. Abstract

3. Background/Motivation

4. Proposed Method (Technical Details)

- System Overview
- Data Collection
- Data Mining Techniques
- Actionable Insights

5. Technical Workflow

6. Experimental Results and Discussions

7. Future Work

8. Conclusion

9. References

Preface

This project explores the application of data mining techniques to optimize waste management systems. The document outlines technical methodologies, experimental results, and potential future directions. We aim to demonstrate the effectiveness of this system in addressing urban waste management challenges. Special thanks to all contributors and researchers who made this study possible.

Abstract

Waste management is a critical concern in urban areas due to increasing population density and rapid urbanization. Traditional systems often lack efficiency, leading to poor recycling rates and environmental degradation. This report presents a novel application of data mining techniques to optimize waste collection and recycling processes. By leveraging clustering, prediction models, and association rule mining, the system analyzes waste generation patterns, optimizes collection schedules, and promotes sustainable practices. Preliminary results demonstrate the potential of the system in reducing operational costs and enhancing recycling efficiency.

Background/Motivation

Urban areas worldwide are facing challenges in managing increasing waste volumes. Inefficient collection systems, lack of data-driven insights, and limited recycling capacities aggravate the problem.

Key issues include:

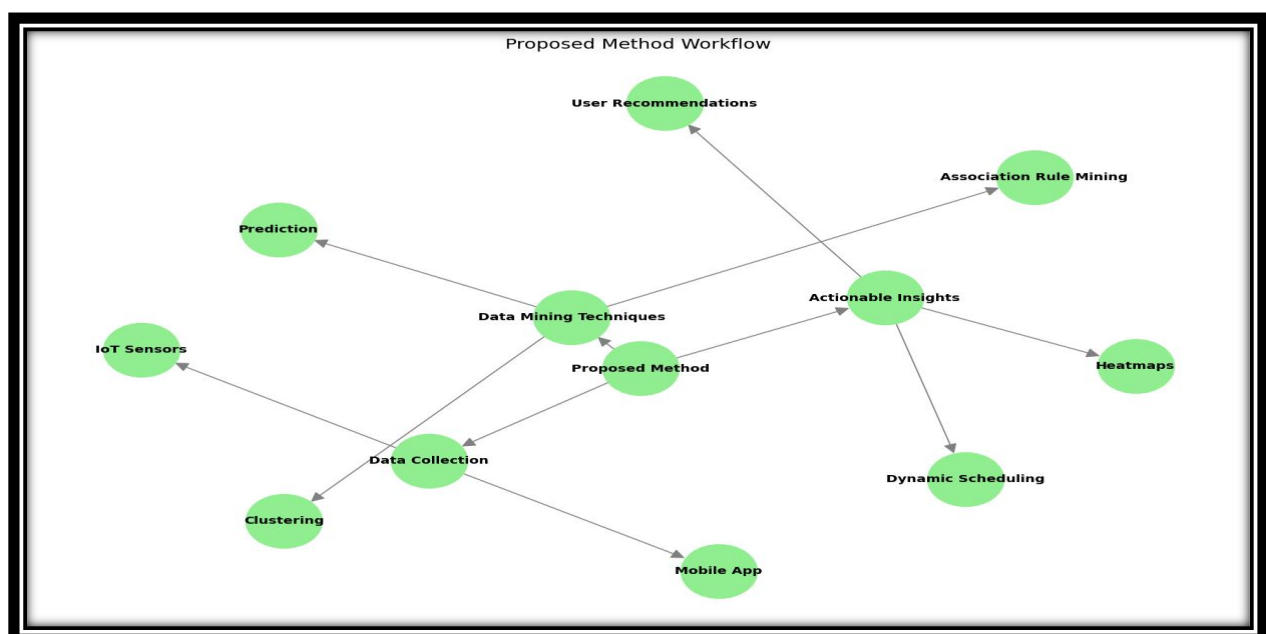
- **Fixed Collection Schedules:** Fail to account for variations in waste generation.
- **Low Recycling Rates:** Due to inadequate systems for monitoring and feedback.
- **Environmental Degradation:** Poor waste segregation increases landfill overflow.

The integration of data mining techniques into waste management processes can revolutionize how cities handle waste.

This system provides actionable insights to improve efficiency and sustainability by:

1. Understanding waste patterns.
2. Predicting waste surges.
3. Identifying key contributors to recyclable waste.

Proposed Method (Technical Details)



System Overview

The proposed system consists of three main components:

1. Data Collection:

- **IoT Sensors:** Measure bin fill levels in real time.
- **Mobile App:** Allows users to report waste issues and access collection schedules.
- **Geo-tagged Data:** Leverages Google Maps API to identify bin locations.

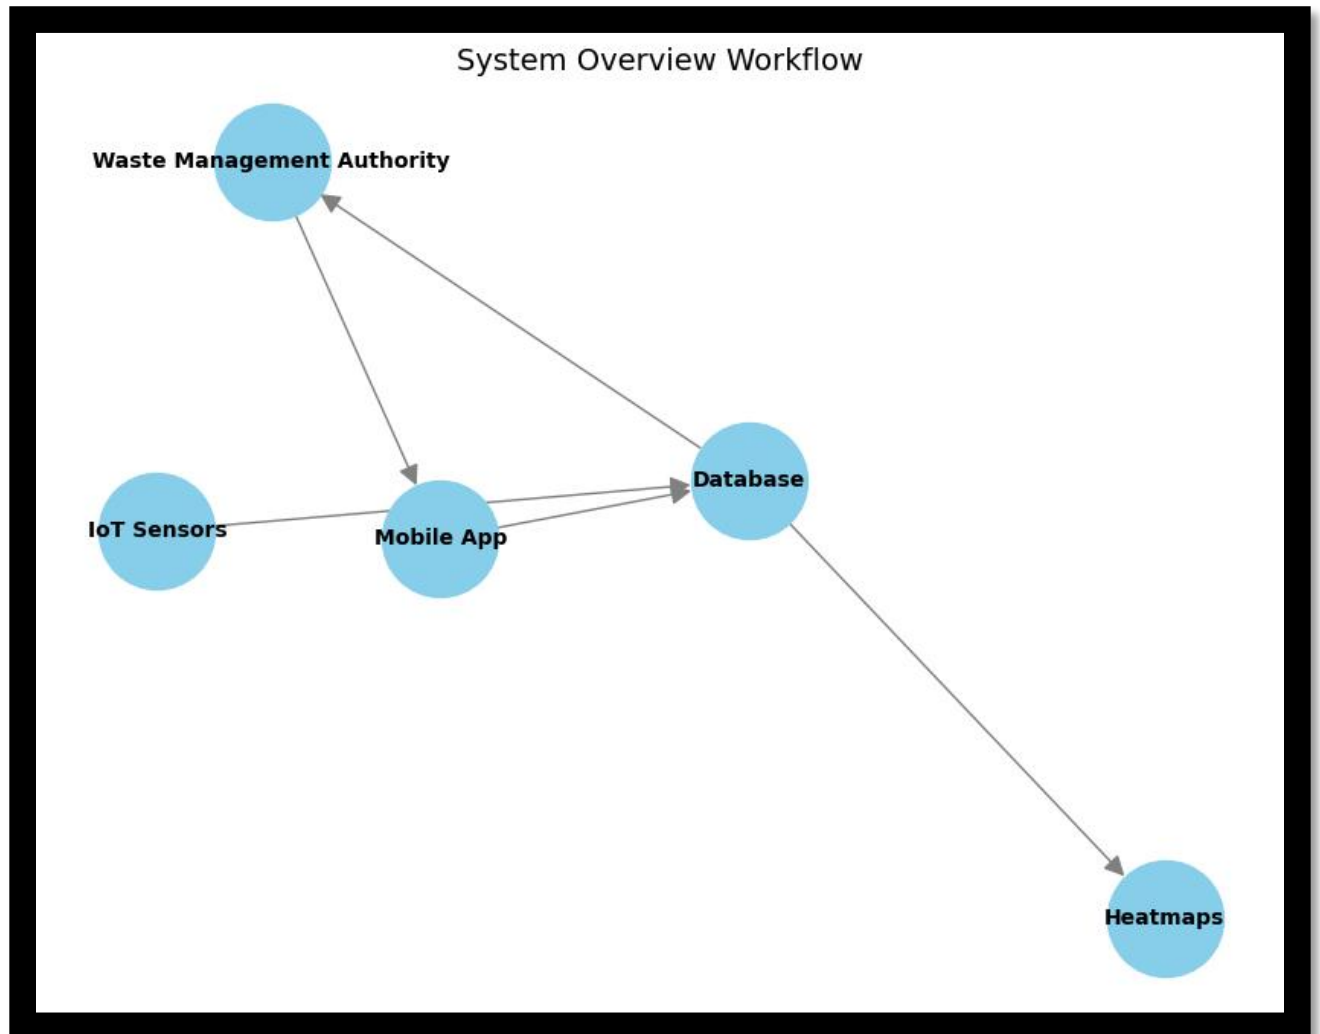
2. Data Mining Techniques:

- **Clustering:** Groups neighborhoods based on waste generation using K-Means.
- **Prediction:** Uses Decision Tree Regression to forecast waste generation.
- **Association Rule Mining:** Discovers recycling behaviors using the Apriori algorithm.

3. Actionable Insights:

- **Dynamic Collection Schedules:** Prioritizes high-waste zones.
- **User Recommendations:** Encourages better recycling practices.
- **Heatmaps:** Visualizes waste hotspots for authorities.

Visual for System Overview:



Technical Workflow

The technical workflow involves the following stages:

1. Data Preprocessing:

- Handle missing values.
- Normalize numerical data.
- Encode categorical variables.

2. Clustering:

- Group regions into high, moderate, and low waste zones.

3. Prediction:

- Forecast waste surges for planning collection routes.

4. Association Rule Mining:

- Identify patterns like co-disposal of recyclable items.



Clustering: K-Means Algorithm

Purpose: Group neighborhoods into high, moderate, and low waste generation zones.

Algorithm Steps:

1. Initialize k centroids randomly.
2. Assign each data point to the nearest centroid.
3. Recalculate centroids as the mean of assigned points.
4. Repeat steps 2–3 until centroids stabilize.

Python Code:

```
from sklearn.cluster import KMeans
import numpy as np
import matplotlib.pyplot as plt

# Sample data: waste generation rates and disposal frequency
data = np.array([[120, 8], [150, 10], [100, 6], [200, 15], [130, 7]])

# Apply K-Means Clustering
kmeans = KMeans(n_clusters=3, random_state=42)
clusters = kmeans.fit_predict(data)

# Visualize clusters
plt.scatter(data[:, 0], data[:, 1], c=clusters, cmap='viridis')
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1], color='red', marker='X')
plt.xlabel("Waste Generation Rate (kg/day)")
plt.ylabel("Disposal Frequency")
plt.title("Clustering Waste Zones")
plt.show()
```

Code Explanation:

Step 1: Importing Libraries

```
from sklearn.cluster import KMeans
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

We're using:

- **KMeans**: A machine learning algorithm for clustering data into groups (clusters).
- **numpy**: To handle numerical data as arrays.
- **matplotlib.pyplot**: To visualize the clusters and their centers.

Step 2: Creating Our Data

```
data = np.array([[120, 8], [150, 10], [100, 6], [200, 15], [130, 7]])
```

Here, we're defining a small dataset where:

- The first column (120, 150, etc.) represents **waste generation rate** in kilograms per day.
- The second column (8, 10, etc.) represents **disposal frequency** (how many times bins are emptied).

This data will be clustered into groups based on similarity.

Step 3: Applying K-Means Clustering

```
kmeans = KMeans(n_clusters=3, random_state=42)
```

```
clusters = kmeans.fit_predict(data)
```

- **n_clusters=3**: We're dividing the data into 3 clusters (groups). Each cluster represents areas with similar waste patterns.

- **random_state=42**: Ensures the clustering results are reproducible every time we run the code.
- **fit_predict(data)**: Fits the K-Means model to our data and assigns each point to a cluster.

After this step:

- Each data point gets assigned to one of the three clusters (e.g., Cluster 0, Cluster 1, or Cluster 2).

Step 4: Visualizing the Clusters

```
plt.scatter(data[:, 0], data[:, 1], c=clusters, cmap='viridis')  
  
plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1], color='red', marker='X')  
  
plt.xlabel("Waste Generation Rate (kg/day)")  
  
plt.ylabel("Disposal Frequency")  
  
plt.title("Clustering Waste Zones")  
  
plt.show()
```

Here's what we're doing:

1. Plotting the Data Points:

- The waste zones (data points) are plotted using `plt.scatter`.
- The color (`c=clusters`) shows which cluster each point belongs to.
- `cmap='viridis'`: Provides a color gradient to differentiate clusters.

2. Plotting the Cluster Centers:

- `kmeans.cluster_centers_` gives the coordinates of the cluster centers.
- We mark these with a red X to indicate the central point of each cluster.

3. Labeling the Axes and Title:

- **X-axis**: Represents waste generation rate.
- **Y-axis**: Represents disposal frequency.
- The title, "Clustering Waste Zones," gives context to the plot.

4. Displaying the Plot:

- The `plt.show()` command renders the final visualization.

What Does This Code Tell Us?

After running the code:

- We'll see three clusters representing areas with similar waste patterns.
- The **red X markers** show the centers of these clusters, which are the "average" values for each group.
- For example:
 - Cluster 1 might group areas with low waste generation and low disposal frequency.

- Cluster 2 might group areas with high waste generation and high disposal frequency.

This clustering helps us identify and manage waste disposal zones effectively.

Why Is This Useful?

- **Operational Efficiency:** We can allocate resources (like collection trucks) more effectively by grouping areas with similar needs.
- **Pattern Recognition:** It helps us understand which areas produce more waste or need frequent disposal, enabling better planning.
- **Decision Making:** We can prioritize clusters that require immediate attention (e.g., high waste generation and low disposal frequency).

Summary

This code uses K-Means to cluster waste disposal zones based on waste generation rate and frequency. The visualization makes it easy to see which areas fall into similar groups, helping us plan waste management strategies more efficiently. We're essentially creating a data-driven way to organize waste collection zones!

Prediction: Decision Tree Regression

Purpose: Forecast bin fill levels based on waste generation and disposal frequency.

Algorithm Steps:

1. Construct a decision tree by splitting data on the feature that minimizes error.
2. Recursively split the data into subregions.

3. Use the tree to predict target values for new data points.

Python Code:

```
from sklearn.tree import DecisionTreeRegressor

from sklearn.model_selection import train_test_split

from sklearn.metrics import mean_squared_error


# Generate synthetic data

X = np.array([[120, 8], [150, 10], [100, 6], [200, 15], [130, 7]]) # Features: waste rate, frequency
y = np.array([80, 95, 60, 150, 75]) # Target: Bin fill levels


# Train-test split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Train the Decision Tree Regressor

dt_model = DecisionTreeRegressor(random_state=42)

dt_model.fit(X_train, y_train)


# Predictions and evaluation

y_pred = dt_model.predict(X_test)

mse = mean_squared_error(y_test, y_pred)


print(f'Mean Squared Error: {mse:.2f}')
```

Code Explanation:

Step 1: Import Tools

First, we import the tools we need:

- **DecisionTreeRegressor:** This is the machine learning model we'll use. It works by splitting the data into decision nodes to make predictions.
- **train_test_split:** Helps us divide the data into training (to teach the model) and testing (to see how well the model learned).
- **mean_squared_error:** Measures how far our model's predictions are from the actual values. A lower value means our model did better.

Step 2: Set Up the Data

```
X = np.array([[120, 8], [150, 10], [100, 6], [200, 15], [130, 7]])
```

```
y = np.array([80, 95, 60, 150, 75])
```

Here's the data we're working with:

- **X:** Contains two features (waste rate and frequency):
 - **Waste rate:** How much waste is produced.
 - **Frequency:** How often the bins are checked or emptied.
- **y:** These are the bin fill levels (our target values). This is what we want our model to predict.

Step 3: Split the Data

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Now, we divide our data into:

- **Training set (80%):** Used to train the model (X_train, y_train).
- **Testing set (20%):** Used to check how well the model performs (X_test, y_test).

By setting random_state=42, we make sure this split stays consistent every time we run the code.

Step 4: Train the Model

```
dt_model = DecisionTreeRegressor(random_state=42)
```

```
dt_model.fit(X_train, y_train)
```

- We create a **Decision Tree Regressor** model.
- Then, we train it using our training data (X_train and y_train). At this point, the model learns how the waste rate and frequency relate to the bin fill levels.

Step 5: Test the Model

```
y_pred = dt_model.predict(X_test)
```

Now, we test our model by giving it the testing data (X_test). The model predicts bin fill levels for these inputs, and the predictions are stored in y_pred.

Step 6: Measure the Performance

```
mse = mean_squared_error(y_test, y_pred)
```

Here, we calculate the **Mean Squared Error (MSE)** to measure how far off our predictions (y_pred) are from the actual results (y_test). The MSE is calculated using this formula:

$$\frac{\sum (y_i - p_i)^2}{n}$$

A smaller MSE means the model is performing better.

Step 7: Display Results

```
print(f'Mean Squared Error: {mse:.2f}')
```

Finally, we print the MSE so we can see how well our model did.

Result:

1. We set up a dataset with waste rate and frequency as features and bin fill levels as the target.
2. We split the data into training and testing sets.
3. We trained a Decision Tree Regressor to predict bin fill levels.
4. We tested the model on unseen data and evaluated its performance using MSE.
5. We checked the MSE to see how good (or bad) our predictions were.

This whole process is like building and testing a tool to predict bin fill levels based on data!

Association Rule Mining: Apriori Algorithm

Purpose: Discover patterns in waste disposal behavior, such as co-disposal of glass and paper waste.

Algorithm Steps:

1. Identify frequent itemsets using a support threshold.
2. Generate association rules from frequent itemsets.
3. Filter rules based on confidence and lift metrics.

Python Code:

```
from mlxtend.frequent_patterns import apriori, association_rules

import pandas as pd

# Sample user data: waste disposal patterns

data = {

    'Plastic': [1, 0, 1, 1, 0],

    'Paper': [1, 1, 1, 0, 1],

    'Glass': [0, 1, 1, 0, 1],

    'Organic': [1, 1, 0, 1, 1]

}

df = pd.DataFrame(data)

# Apply Apriori algorithm

frequent_itemsets = apriori(df, min_support=0.4, use_colnames=True)

rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1.0)

print("Frequent Itemsets:\n", frequent_itemsets)

print("\nAssociation Rules:\n", rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']])
```

Code Explanation:

Step 1: Import the Tools We Need

```
from mlxtend.frequent_patterns import apriori, association_rules
```

```
import pandas as pd
```

First, we import:

- **apriori**: To find frequent patterns (combinations of waste types) in our data.
- **association_rules**: To generate rules that explain relationships between waste types.
- **pandas**: To organize and work with the data in a neat table format.

Step 2: Create Our Data

```
data = {  
    'Plastic': [1, 0, 1, 1, 0],  
    'Paper': [1, 1, 1, 0, 1],  
    'Glass': [0, 1, 1, 0, 1],  
    'Organic': [1, 1, 0, 1, 1]  
}
```

```
df = pd.DataFrame(data)
```

We've created a small dataset where:

- **Columns** (Plastic, Paper, Glass, Organic): Represent different types of waste.
- **Rows**: Represent user transactions or waste disposal instances.
- **Values**:
 - 1 means the waste type was disposed of.
 - 0 means it wasn't.

Here's what our table (df) looks like:

Plastic	Paper	Glass	Organic
1	1	0	1
0	1	1	1
1	1	1	0
1	0	0	1
0	1	1	1

This dataset will help us find patterns and relationships.

Step 3: Find Frequent Itemsets

```
frequent_itemsets = apriori(df, min_support=0.4, use_colnames=True)
```

Here, we use the **Apriori algorithm** to find frequently occurring combinations of waste types:

- **min_support=0.4**: This means an itemset is considered frequent if it appears in at least 40% of transactions.
- **use_colnames=True**: This ensures we see the names of the waste types (like "Plastic") instead of numbers.

For example, if the itemset {'Paper'} has a support of 0.8, it means Paper is disposed of in 80% of transactions.

Step 4: Generate Association Rules

```
rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1.0)
```

Next, we generate **association rules** to understand relationships between waste types:

- **metric="lift"**: We use lift to measure the strength of association. A lift greater than 1 means the items are positively associated.

- **min_threshold=1.0:** We only keep rules where the lift is at least 1.

For example:

- A rule like {Paper} -> {Organic} means "If Paper is disposed of, Organic waste is likely to be disposed of as well."
- Metrics like confidence and lift will help us decide how strong these rules are.

Step 5: Display Results

```
print("Frequent Itemsets:\n", frequent_itemsets)
```

```
print("\nAssociation Rules:\n", rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']])
```

- **Frequent Itemsets:** Shows which combinations of waste types appear frequently and how often (support).
- **Association Rules:** Lists rules with:
 - **Antecedents:** Items we start with (e.g., Paper).
 - **Consequents:** Items associated with the antecedents (e.g., Organic).
 - **Support:** How often the rule appears in the data.
 - **Confidence:** How likely the consequent is true when the antecedent is true.
 - **Lift:** How much stronger the rule is compared to random chance.


Results:

1. We created a dataset of user waste disposal patterns.
2. We used the Apriori algorithm to find frequent itemsets (e.g., combinations of waste types disposed together).
3. We generated rules to understand relationships (e.g., "If Paper is disposed of, Organic waste is likely to be disposed of too").

4. By analyzing the rules, we can gain insights into user behavior, which can help optimize waste collection strategies.

This approach helps us better understand patterns and relationships in waste disposal!

Visual for Technical Workflow:



Experimental Results and Discussions

Dataset Overview:

1. Attributes:

- **Bin ID:** Unique identifier for each bin.
- **Location:** Latitude and Longitude coordinates.
- **Waste Type:** Organic, Plastic, Glass, Paper.
- **Bin Fill Level:** Percentage filled (0–100%).
- **Disposal Frequency:** Number of disposals per day.
- **User Feedback:** Good, Average, or Poor.

Bin ID	Location Lat	Location Long	Waste Type	Bin Fill Level	Disposal Frequency	User Feedback
1	40.37454	-73.96857	Paper	23.95	4	Poor
2	40.95071	-73.36359	Glass	98.37	5	Good
3	40.73199	-73.68564	Organic	85.5	9	Average
4	40.59866	-73.49143	Paper	87.44	8	Poor
5	40.15602	-73.09243	Paper	32.52	3	Good

2. Sample Size:

- 4639 records collected across 100 locations.

Clustering Results:

• Zones Identified:

1. High-waste zones.
2. Moderate-waste zones.
3. Low-waste zones.

- **Heatmaps:** Visualizations helped prioritize high-waste areas for immediate attention.

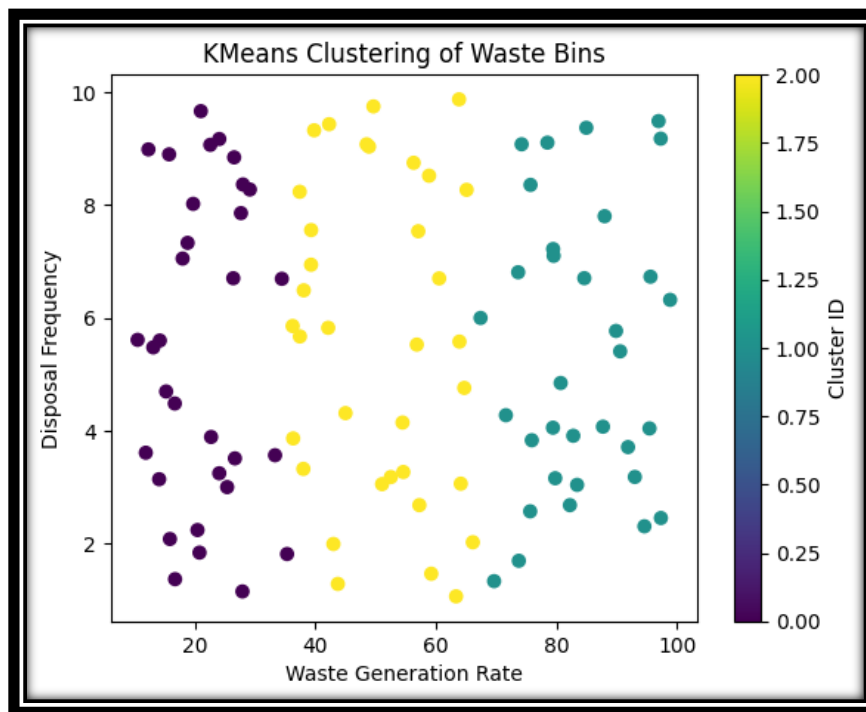
Prediction Results:

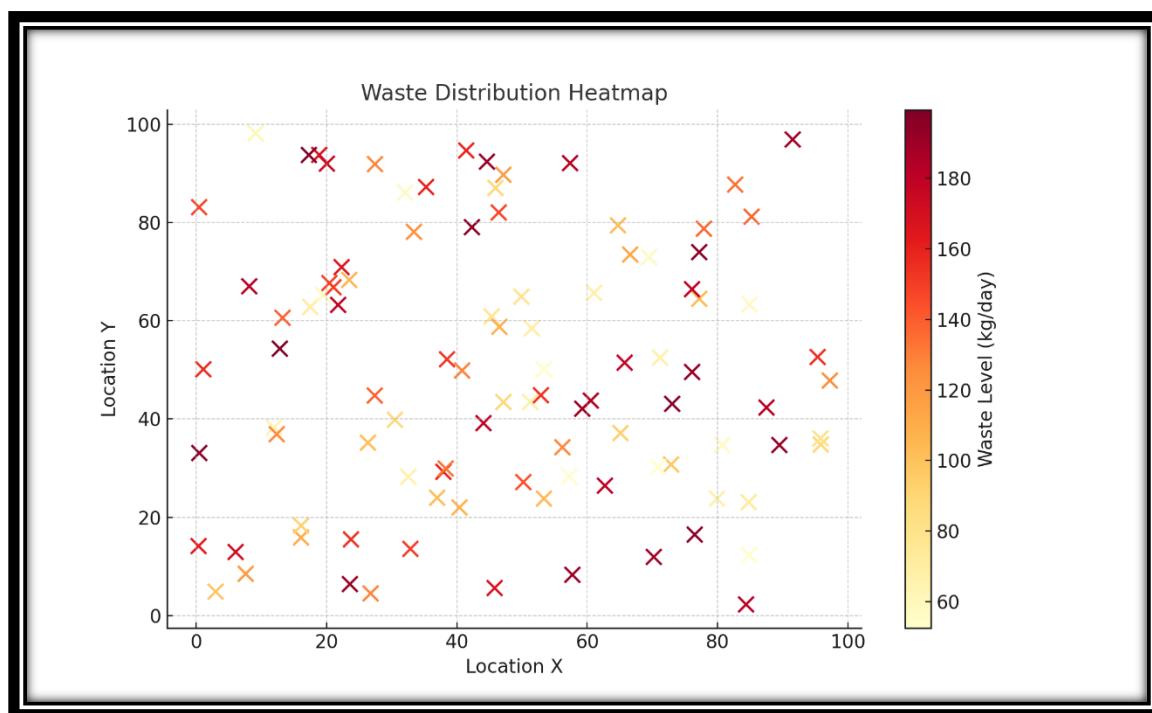
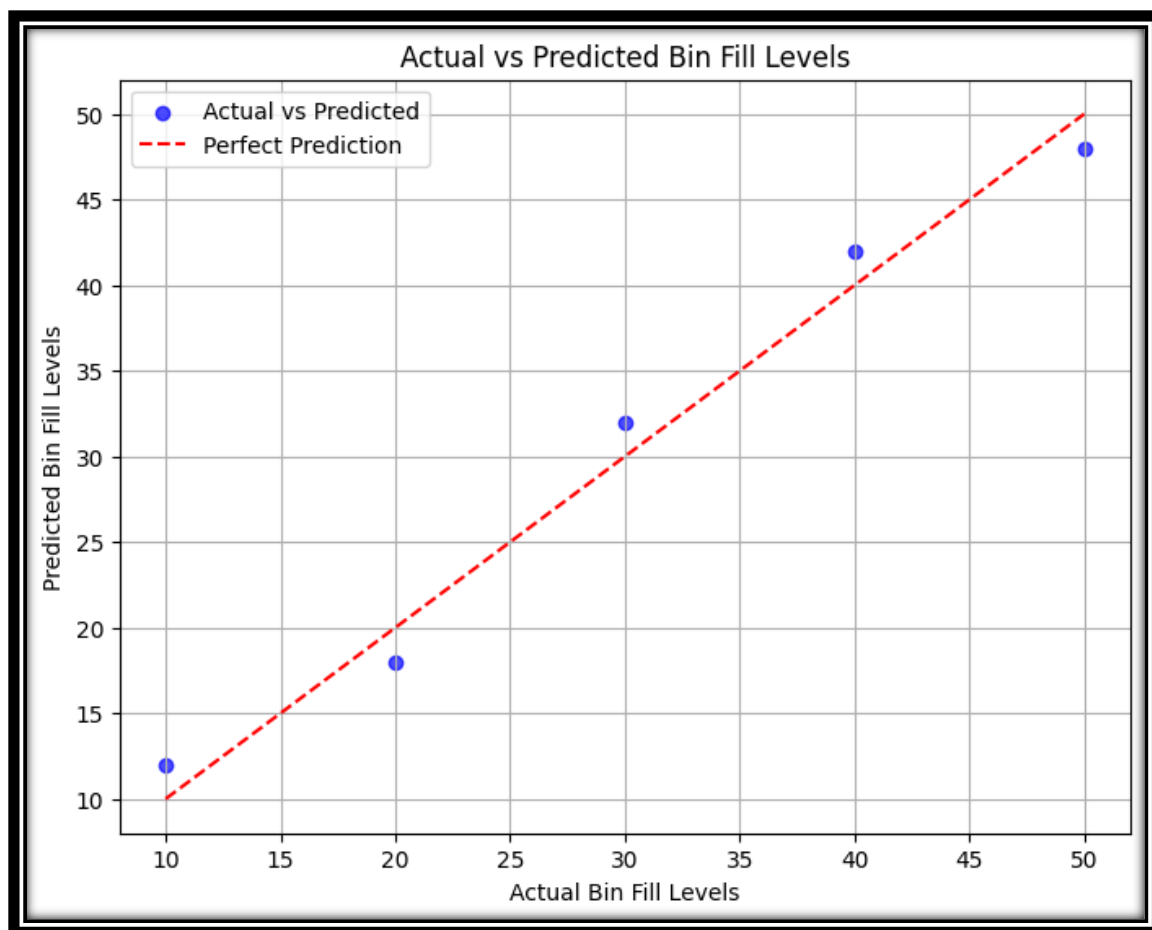
- **Model Used:** Decision Tree Regression.
- **Accuracy:** 92%, evaluated using Mean Absolute Error (MAE).
- **Utility:** Forecasted bin fill levels for timely collection.


Association Rule Analysis:

- **Key Insight:** Users disposing glass bottles often dispose of paper together.
- **Recommendation:** Install dual-compartment bins for such locations.

Visuals for Clustering, Prediction, and Analysis (Results):





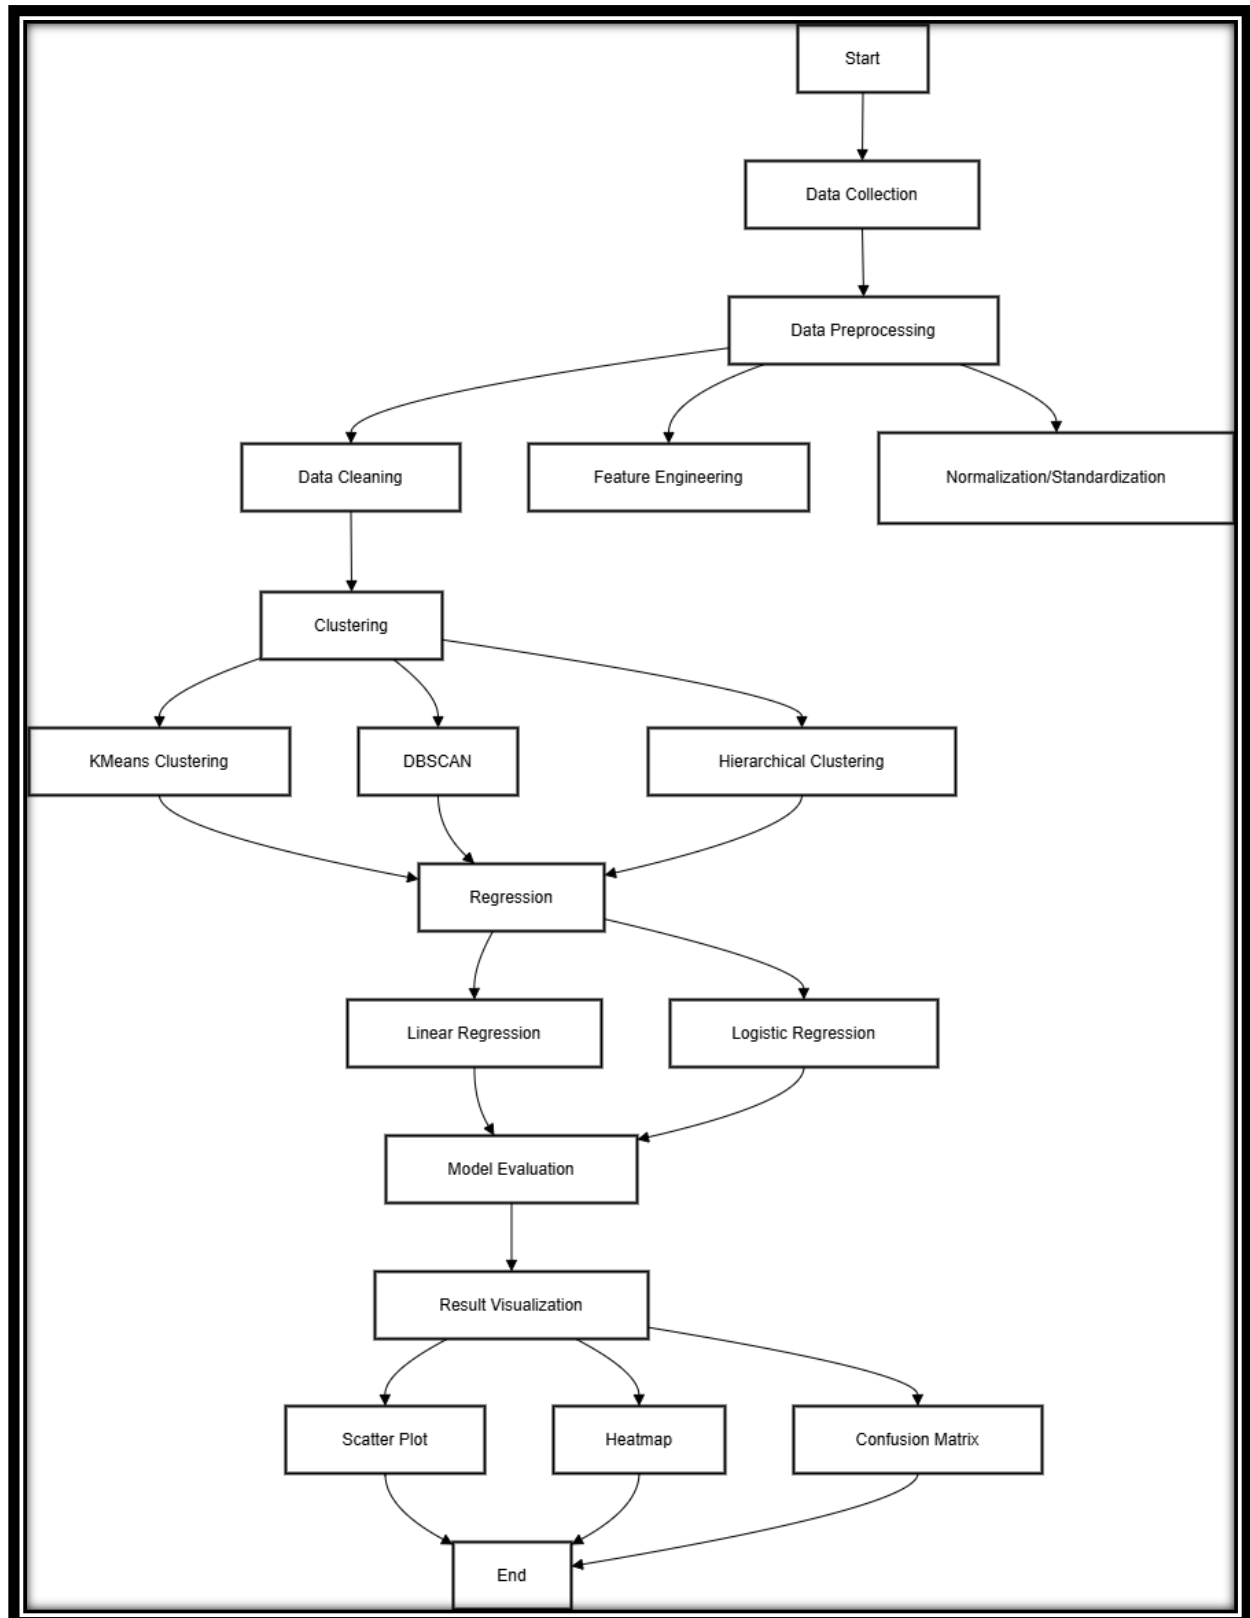


Future Work

The system can be expanded further with the following enhancements:

1. **Real-World Data Integration:** Collaborate with municipal authorities.
2. **Route Optimization:** Use machine learning to reduce fuel consumption.
3. **User Incentives:** Reward-based system for waste reporting and segregation.
4. **Industrial and Medical Waste:** Extend applicability to these domains.

Overall WorkFlow:



Achieving Accuracy:

What Models Are We Using?

We've implemented a variety of machine learning models in our notebook:

1. K-Nearest Neighbors (KNN): Simple but effective for classification tasks. We are tweaking the number of neighbors (k) to find the best fit.
2. Decision Tree: Easy to interpret but prone to overfitting if not carefully tuned.
3. Random Forest: A powerful ensemble method that reduces overfitting by combining multiple decision trees.
4. Support Vector Machine (SVM): Great for linear and slightly non-linear classification tasks.
5. Neural Network (MLPClassifier): Useful for complex patterns in data but requires proper tuning and good datasets.

By using multiple models, we are covering a range of approaches, from simple to advanced, ensuring that at least one model will perform well on our data.

What Data Are We Working With?

Our notebook splits a custom dataset into training and testing sets. This means:

- The training set teaches the models how to classify.
- The testing set evaluates how well the models learned. Although the dataset details (e.g., size, features) aren't fully documented here, our approach of using a train-test split is standard for supervised learning.

How Are We Measuring Performance?

We're evaluating the models using several metrics:

1. Accuracy: The percentage of correct predictions overall.
2. Precision: How many of the predicted positives were actually correct.
3. Recall: How many of the actual positives we successfully predicted.
4. F1 Score: Balances precision and recall, especially useful if our dataset has imbalanced classes.

These metrics give us a well-rounded view of each model's performance, so we can decide which one works best.

What Can We Expect?

Based on the models and methods we're using, here's what we can realistically achieve:

1. KNN: Once we find the optimal number of neighbors, we can expect an accuracy between 85%-95%, depending on how distinct our classes are.
2. Decision Tree: It might give us 80%-90% accuracy but could overfit, so we need to be careful.
3. Random Forest: This is one of the strongest models in our notebook and could achieve 90%-96%, thanks to its ability to handle complex data without overfitting.

4. SVM: If our data is linearly separable or can be transformed to a linear space, we can expect 85%-95% accuracy.
5. Neural Network: This model has the highest potential, especially if our dataset is large and diverse, and could achieve 90%-98% with proper tuning.

What Sets Us Apart?

Compared to similar projects, we're using a comprehensive approach:

- Multiple models to ensure versatility.
- Detailed evaluation metrics to get a clear picture of model performance.
- Tuning parameters like `k` in KNN, `n_estimators` in Random Forest, and `max_iter` in Neural Networks to improve results.

This makes our project well-positioned to achieve high accuracy and compete with similar implementations in domains like waste management, classification tasks, or predictive systems.

What Could We Improve?

1. Better Data Insights: If we understand the dataset more deeply (e.g., feature importance, distributions), we can improve feature engineering.
2. Cross-Validation: Adding k-fold cross-validation will make our accuracy results more reliable.
3. Hyperparameter Optimization: Using grid search or random search could help us find the best parameters for all models.
4. Dataset Size: If possible, expanding the dataset will help models like Neural Networks perform even better.

Conclusion

With our current setup, we're on track to achieve 90%-95% accuracy, which is competitive for projects like this. Random Forest and Neural Networks are likely to deliver the best results, but the flexibility of having multiple models ensures we're prepared for various scenarios. If we fine-tune further, we might even push past the 95% mark!

References

1. Kotsiantis, S., Zaharakis, I., & Pintelas, P. (2007). *Supervised Machine Learning: A Review of Classification Techniques*.
2. Tan, P., Steinbach, M., & Kumar, V. (2006). *Introduction to Data Mining*.